

CCDS Workshop Task Results

Submitted By: Syed Md. Ahnaf Hasan

1. TODO Task-1:

Architecture-1:

The custom CNN model was trained on the train set. During training the model was also simultaneously evaluated on the validation set. Later, the test split was used to evaluate the performance of the trained model.




Fig: Loss curve




Fig: Confusion Matrix of Test Split

	precision	recall	f1-score	support
airplane	0.79	0.87	0.83	567
automobile	0.88	0.90	0.89	611
bird	0.71	0.70	0.70	623
cat	0.59	0.62	0.61	566
deer	0.72	0.77	0.75	585
dog	0.73	0.70	0.71	609
frog	0.84	0.79	0.81	597
horse	0.86	0.80	0.83	622
ship	0.89	0.87	0.88	606
truck	0.88	0.85	0.87	614
accuracy			0.79	6000
macro avg	0.79	0.79	0.79	6000
weighted avg	0.79	0.79	0.79	6000

Fig: Performance metrics on the Test set

Observations:

- Overfitting was noticed after epoch 2 in training.
- The model's performance was best for the automobile class (548 correct classifications) and worst for the cat class (352 correct classifications) as seen from the confusion matrix of the test set.
- The macro average F1 score is 0.79.

2. TODO Task-2:

Architecture-2:

- Changes in code:
Sigmoid layers were used instead of ReLU layers after each linear layers:

```
nn.Flatten(),
nn.Linear(256*4*4, 1024),
nn.Sigmoid(),
nn.Linear(1024, 512),
nn.Sigmoid(),
nn.Linear(512, output_size))
```

Architecture-3:

- Changes in code:
An additional Conv2d layer and a ReLU layer was added after the last maxpool layer before flattening:

```
nn.MaxPool2d(2, 2), # output: 256 x 4 x 4
nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1), # output: 256 x 4
nn.ReLU(),

nn.Flatten(),
```

Architecture-4:

- Changes in code:
Batch normalization was used after each linear layer before the Sigmoid activation functions:

```
nn.Linear(256*4*4, 1024),
nn.BatchNorm1d(num_features = 1024),
nn.Sigmoid(),
nn.Linear(1024, 512),
nn.BatchNorm1d(num_features = 512),
nn.Sigmoid(),
```

Loss Curves:




Fig: Loss Curve (Architecture-2)




Fig: Loss Curve (Architecture-3)




Fig: Loss Curves (Architecture-4)

Observation:

- Introducing Sigmoid activations after linear layers in architecture-2 makes the gradients of the high activation outputs close to zero, consequently the weights are not updated in the backpropagation and the loss does not decrease.
- Introducing the extra Conv+Relu block in architecture -3 makes the loss drop after 6 epochs.
- Batch normalization in architecture-4 normalizes the activation outputs of the linear layers. This forces the activation outputs to stay in a linear range where gradients are not zero and hence the updates take place during backpropagation. However overfitting is noticed after epoch 4.

3. TODO Task-3:

Comparison table of the 4 architectures is shown below:

Table: Comparing Performance Metrics

Metrics Models	Accuracy	Precision (Macro Avg)	Recall (Macro Avg)	F1 Score (Macro Avg)
Architecture-1	0.78	0.78	0.78	0.78
Architecture-2	0.22	0.19	0.22	0.17
Architecture-3	0.43	0.43	0.43	0.42
Architecture-4	0.82	0.82	0.82	0.82

Observation:

The architecture-4 obtained the highest scores across all the performance metrics.

4. TODO Task-4:

I have used architecture-4 for this task, since I have obtained the highest accuracy by using this architecture.

Changes in code:

For Xavier initialization the `_initialize_weights_()` method was declared which was called in the `__init__()` method.

```
nn.Dropout(p=0.5),
nn.Linear(512, output_size))
self._initialize_weights_()

def _initialize_weights(self):
    for layer in self.children():
        if isinstance(layer, nn.Linear):
            nn.init.xavier_uniform_(layer.weight)
            nn.init.zeros_(layer.bias)
```

For dropout the nn.Dropout layer was added after the activation functions:

```
nn.Linear(256*4*4, 1024),
nn.BatchNorm1d(num_features = 1024),
nn.Sigmoid(),
nn.Dropout(p=0.5),
nn.Linear(1024, 512),
nn.BatchNorm1d(num_features = 512),
nn.Sigmoid(),
nn.Dropout(p=0.5),
```

The learning_rate scheduler was changed to cosine learning rate:

```
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max = 10)
```

The loss curves for batch size = 256, learning rate = 0.001, and Cosine Annealing Learning rate scheduler are shown below:




Fig: Loss Curves




Fig: Confusion matrix

	precision	recall	f1-score	support
airplane	0.81	0.86	0.84	584
automobile	0.94	0.91	0.92	578
bird	0.77	0.69	0.73	595
cat	0.62	0.70	0.66	584
deer	0.78	0.77	0.77	589
dog	0.74	0.76	0.75	634
frog	0.88	0.83	0.85	618
horse	0.88	0.85	0.87	624
ship	0.90	0.90	0.90	589
truck	0.87	0.89	0.88	605
accuracy			0.82	6000
macro avg	0.82	0.82	0.82	6000
weighted avg	0.82	0.82	0.82	6000

Fig: Performance Metrics

Observation:

The performance after making the aforementioned changes is similar to the original architecture's. Overfitting is noticeable in this configuration also.

5. TODO Task-5:

The confusion matrices for torch embeddings and word2vec embeddings are shown below:




Fig: Conf. Matrix (torch Embeddings)

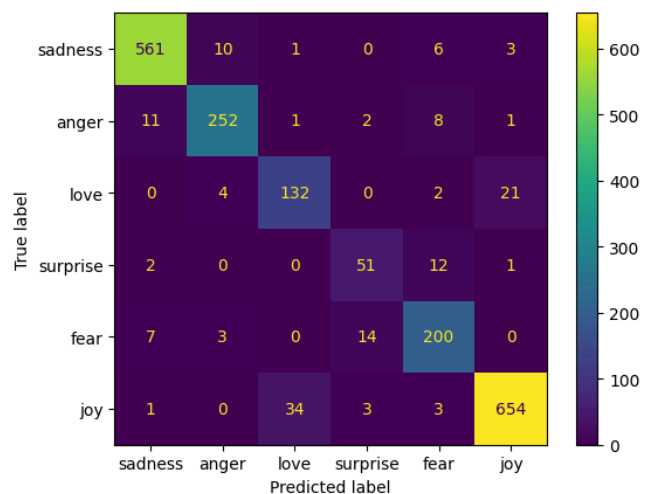


Fig: Conf. Matrix (Word2Vec Embeddings)

	precision	recall	f1-score	support
sadness	0.95	0.98	0.97	581
anger	0.96	0.87	0.92	275
love	0.80	0.77	0.79	159
surprise	0.77	0.77	0.77	66
fear	0.86	0.90	0.88	224
joy	0.94	0.94	0.94	695
accuracy			0.92	2000
macro avg	0.88	0.87	0.88	2000
weighted avg	0.92	0.92	0.92	2000

Fig: Performance metrics (torch emb.)

	precision	recall	f1-score	support
sadness	0.96	0.97	0.96	581
anger	0.94	0.92	0.93	275
love	0.79	0.83	0.81	159
surprise	0.73	0.77	0.75	66
fear	0.87	0.89	0.88	224
joy	0.96	0.94	0.95	695
accuracy			0.93	2000
macro avg	0.87	0.89	0.88	2000
weighted avg	0.93	0.93	0.93	2000

Fig: Performance metrics (Word2Vec)

- **Observation:**

The macro average F1 scores for both torch embeddings and word2vec embeddings was found to be 0.88. However, the accuracy was slightly higher in case of word2vec embeddings.

6. TODO Task-6:

- **Pseudo plan:**

1. Observe the attention output's shape and its elements to understand its structure. -> The shape of the attention output is: (batch x encoder_block_num x num_of_sentences x num_of_attention_heads_per_encoder_block x num_of_tokens x num_of_tokens).
2. Try to find the most important token for a single sentence.
3. Create a function to output the predicted class, most important token, highest attention score for a single sentence.
4. Iterate through each sentence in each batch and aggregate the attention scores for each tokens in each class.
5. Display the most important token in each class.

- **Colab notebook link:**

https://colab.research.google.com/drive/1iNXyahNct_WOt-SJxarZAJVizeTmEurs?usp=sharing

- **Observation:**

The most attentive tokens found for each class are as follows:

1. Sadness -> depressed
2. Anger -> feeling
3. Love -> fond
4. Surprise -> overwhelmed
5. Fear -> shy
6. Joy -> feel